

```

!pip install --quiet wandb
!pip install --quiet transformers
!pip install --quiet datasets
!pip install --quiet emoji
!pip install --quiet kaggle
!pip install --quiet torchinfo
!pip install --quiet imbalanced-learn
!pip install --quiet gdown
!pip install --quiet clean-text
!pip install --quiet accelerate -U
!pip install --quiet transformers[torch]
!pip install --quiet huggingface_hub[tensorflow]
!pip install --quiet huggingface_hub[cli,torch]
!pip install --quiet openai

from huggingface_hub import HfApi
import huggingface_hub

#transformers
import transformers
from transformers import AutoTokenizer
from transformers import GPT2Tokenizer
from transformers import BertTokenizer
from transformers import OpenAIGPTTokenizer
from transformers import AutoConfig
from transformers import GPT2Config
from transformers.models.bert.modeling_bert import BertModel
from transformers.models.gpt2.modeling_gpt2 import GPT2Model
from transformers.models.openai.modeling_openai import OpenAIGPTModel
from transformers.models.bert.modeling_bert import BertPreTrainedModel
from transformers.models.gpt2.modeling_gpt2 import GPT2PreTrainedModel
from transformers.models.openai.modeling_openai import OpenAIGPTPreTrainedModel
from transformers import BertForSequenceClassification
from transformers import GPT2ForSequenceClassification
from transformers import OpenAIGPTForSequenceClassification
from transformers.modeling_outputs import SequenceClassifierOutput
from transformers import Trainer
from transformers import TrainingArguments
from transformers import get_scheduler
from transformers import TextDataset
from transformers.modeling_utils import PreTrainedModel

#torch
import torch
import torch.nn as nn
from torch.nn import BCEWithLogitsLoss, MSELoss, CrossEntropyLoss
from torch.optim import AdamW
from torch.utils import data
from torch.utils.data import Dataset as ds, DataLoader

# dataset
import datasets
from datasets import Dataset
from datasets import Sequence
from datasets import Value
from datasets import Features
from datasets import ClassLabel
from datasets import DatasetDict

# sklearn
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import ConfusionMatrixDisplay, confusion_matrix

#open ai
import openai

# others
import os
import wandb
import re, string
import emoji
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

```

```
import numpy as np
import pickle
import joblib
import traceback
from tqdm import tqdm
from collections import defaultdict
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
api = HfApi()
```

```

===== 2.1/2.1 MB 29.1 MB/s eta 0:00:00
===== 188.6/188.6 kB 23.3 MB/s eta 0:00:00
===== 218.8/218.8 kB 24.5 MB/s eta 0:00:00
Preparing metadata (setup.py) ... done
===== 62.7/62.7 kB 7.0 MB/s eta 0:00:00
Building wheel for pathtools (setup.py) ... done
===== 7.6/7.6 MB 24.3 MB/s eta 0:00:00
===== 268.8/268.8 kB 27.0 MB/s eta 0:00:00
===== 7.8/7.8 MB 42.9 MB/s eta 0:00:00
===== 1.3/1.3 MB 37.5 MB/s eta 0:00:00
===== 519.6/519.6 kB 2.0 MB/s eta 0:00:00
===== 115.3/115.3 kB 14.8 MB/s eta 0:00:00
===== 194.1/194.1 kB 24.8 MB/s eta 0:00:00
===== 134.8/134.8 kB 18.1 MB/s eta 0:00:00
===== 358.9/358.9 kB 7.0 MB/s eta 0:00:00
===== 175.4/175.4 kB 2.2 MB/s eta 0:00:00
Preparing metadata (setup.py) ... done
===== 53.1/53.1 kB 4.2 MB/s eta 0:00:00
Building wheel for emoji (setup.py) ... done
===== 251.2/251.2 kB 3.4 MB/s eta 0:00:00
===== 67.7/67.7 kB 1.6 MB/s eta 0:00:00
===== 76.5/76.5 kB 1.3 MB/s eta 0:00:00
```

```
! huggingface-cli login --token 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
```

Token will not be saved to git credential helper. Pass `add_to_git_credential=True` if you want to set the git credential as well.
 Token is valid (permission: write).
 Your token has been saved to /root/.cache/huggingface/token
 Login successful

```
huggingface_hub.login()
```



Copy a token from [your Hugging Face tokens page](#) and paste it below.

Immediately click login after copying your token or it might be stored in plain text in this notebook file.

Token:

☒ Add token as git credential?

Login

Pro Tip: If you don't already have one, you can create a dedicated 'notebooks' token with

```

class BertForClassification(BertForSequenceClassification):

    def __init__(self, config):
        super().__init__(config)
        self.num_labels = config.num_labels
        self.config = config

        # Load model body > return all of the HS
        self.bert = BertModel(config)
        # Set up token classification head
        self.dropout = nn.Dropout(config.hidden_dropout_prob)
        self.classifier = nn.Linear(config.hidden_size, config.num_labels)

        # Initialize weights and apply final processing
        self.post_init()

    def forward(self, input_ids=None, attention_mask=None, token_type_ids=None,
                labels=None, **kwargs):
        # Use model body to get encoder representations
        outputs = self.bert(input_ids, attention_mask=attention_mask,
                             token_type_ids=token_type_ids, **kwargs)

        # Apply classifier to encoder representation > [cls]
        sequence_output = self.dropout(outputs[1])
        logits = self.classifier(sequence_output)

        # Calculate losses
        loss = None
        if labels is not None:
            loss_fct = nn.CrossEntropyLoss()
            loss = loss_fct(logits.view(-1, self.num_labels), labels.view(-1))
            # outputs = (loss,) + outputs # can comment

        # return outputs # (loss), logits, (hidden_states), (attentions)

        # Return model output object
        return SequenceClassifierOutput(
            loss=loss,
            logits=logits,
            hidden_states=outputs.hidden_states,
            attentions=outputs.attentions,
        )

model_names = ["bert-large-uncased"]

id2label = {0: 'non-toxic', 1: 'mild', 2: 'toxic'}
label2id = {v:k for (k,v) in id2label.items()}
labels = ['non-toxic', 'mild', 'toxic']

tokenizers = {"bert-large-uncased": BertTokenizer.from_pretrained("bert-large-uncased")}
configs = {"bert-large-uncased": AutoConfig.from_pretrained("bert-large-uncased", num_labels=3, id2label=id2label, label2id=label2id)}
models = {"bert-large-uncased": BertForClassification.from_pretrained("bert-large-uncased", config=configs["bert-large-uncased"])}

Downloading (...)solve/main/vocab.txt: 100%                232k/232k [00:00<00:00, 9.22MB/s]

Downloading (...)okenizer_config.json: 100%                28.0/28.0 [00:00<00:00, 1.96kB/s]

Downloading (...)lve/main/config.json: 100%                571/571 [00:00<00:00, 34.9kB/s]

Downloading model.safetensors: 100%                        1.34G/1.34G [00:05<00:00, 146MB/s]

Some weights of BertForClassification were not initialized from the model checkpoint at bert-large-uncased and are newly initialized: ['
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference

```

```

def clean_text(text):
    text = re.sub('#', '', text) # Removing '#' hashtag
    text = re.sub('\w+:\/\/\S+', '', text) # Removing hyperlink
    text = re.sub('[^a-zA-Z]', ' ', text) # Remove punctuation
    return text.lower()

def compute_metrics(pred):
    labels = pred.label_ids
    preds = pred.predictions.argmax(-1)
    f1 = f1_score(labels, preds, average="weighted")
    acc = accuracy_score(labels, preds)
    return {"accuracy": acc, "f1": f1}

def create_dataset(df, text, label):
    class_names = ['non-toxic', 'mild', 'toxic']
    data_dict = {'text': df[text], 'labels': df[label]}
    tags = ClassLabel(num_classes=3, names=class_names)
    feature_set = Features({'text': Value(dtype='string'), 'labels': tags})

    return Dataset.from_dict(mapping = data_dict, features = feature_set)

def get_dataset():
    train = pd.read_csv('/content/train.csv')
    validation = pd.read_csv('/content/validation.csv')

    dataset_train = create_dataset(train, "message", "target")
    dataset_val = create_dataset(validation, "message", "target")

    train_convert = train.copy()
    train_convert['target'] = train_convert['target'].astype(str)
    train_convert['message'] = (train_convert['message'] + ' ->')
    train_convert['target'] = (' ' + train_convert['target'])
    train_convert.rename(columns={'message': 'prompt', 'target': 'completion'}, inplace=True)
    train_convert.to_json("train_pandas.jsonl", orient='records', lines=True)

    validation_convert = validation.copy()
    validation_convert['target'] = validation_convert['target'].astype(str)
    validation_convert['message'] = (validation_convert['message'] + ' ->')
    validation_convert['target'] = (' ' + validation_convert['target'])
    validation_convert.rename(columns={'message': 'prompt', 'target': 'completion'}, inplace=True)
    validation_convert.to_json("validation_pandas.jsonl", orient='records', lines=True)

    dataset = DatasetDict()
    dataset["train"] = dataset_train
    dataset["validation"] = dataset_val
    return dataset

! wandb login eed5796f17cdf020038ca37377edaaec8c9ddfbe

wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc

wandb.login()
wandb.init(project="bert-toxic-chat-detection", entity="dani-squad")

```

wandb: Currently logged in as: **dfesalbon** (**dani-squad**). Use `wandb login --relogin` to force relogin
Tracking run with wandb version 0.15.10
Run data is saved locally in /content/wandb/run-20230907_001043-rj53tmy9
Syncing run **fancy-blaze-32** to [Weights & Biases \(docs\)](#)
View project at <https://wandb.ai/dani-squad/bert-toxic-chat-detection>
View run at <https://wandb.ai/dani-squad/bert-toxic-chat-detection/runs/rj53tmy9>

```
def finetune_bert(model_name, dataset, max_len=64, train_batch_size=16, num_epochs=5):

    tokenizer = tokenizers[model_name]
    model = models[model_name]

    tokenize = lambda batch: tokenizer(batch["text"], padding=True, truncation=True, max_length=max_len)

    encoded = dataset.map(tokenize, batched=True, batch_size=32)

    model.to(device)
    optimizer = AdamW(model.parameters(), lr=2e-5, eps=1e-8)
    logging_steps = len(encoded['train']) // train_batch_size
    num_training_steps = num_epochs * logging_steps

    scheduler = get_scheduler(
        name="linear",
        optimizer=optimizer,
        num_warmup_steps=0,
        num_training_steps=num_training_steps)

    training_args = TrainingArguments(
        output_dir=f"{model_name}-toxic-detector",
        num_train_epochs=num_epochs,
        per_device_train_batch_size=train_batch_size,
        per_device_eval_batch_size=train_batch_size,
        weight_decay=0.01,
        evaluation_strategy="epoch",
        save_steps=1e6,
        disable_tqdm=False,
        logging_steps=logging_steps,
        push_to_hub=False,
        log_level="error",
        report_to="wandb",
        run_name=model_name)

    optimizers = (optimizer, scheduler)

    model_trainer = Trainer(
def save_model(model_name, model_trainer):
    model_trainer.save_model(f'/content/{model_name}-saved')
    api.upload_folder(
        folder_path=f"/content/{model_name}-saved",
        repo_id=f"dffesalbon/{model_name}-dota-toxic",
        repo_type="model")

dataset = get_dataset()
model_trainer.train()
```

✓ BERT (Large-uncased)

```
model_name, model_trainer, encoded = finetune_bert(model_names[0], dataset)
```

Map: 100% 1722/1722 [00:00<00:00, 2067.87 examples/s]

Map: 100% 192/192 [00:00<00:00, 1157.71 examples/s]

[540/540 04:01, Epoch 5/5]

Epoch	Training Loss	Validation Loss	Accuracy	F1
1	0.857000	0.568056	0.786458	0.781023
2	0.555000	0.417159	0.838542	0.837147
3	0.384100	0.504736	0.796875	0.788389
4	0.231400	0.592420	0.807292	0.803610
5	0.141100	0.615863	0.817708	0.819774

```
try:
    save_model(model_name, model_trainer)
except Exception as ex:
```